My 1st Machine Learning model

Atelier « Maker »

Who we are

Nicolas FROT
Data Squad

Florian BERGAMASCO
EP/EXPLO/GTS/IGR

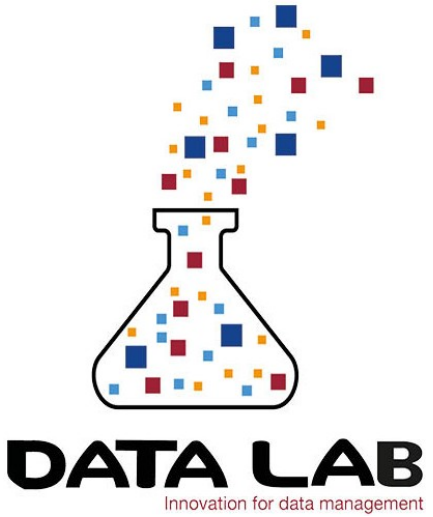

What is Big Data?

1956: 5 Mo,
$50k

2018: 256 Go,
$30

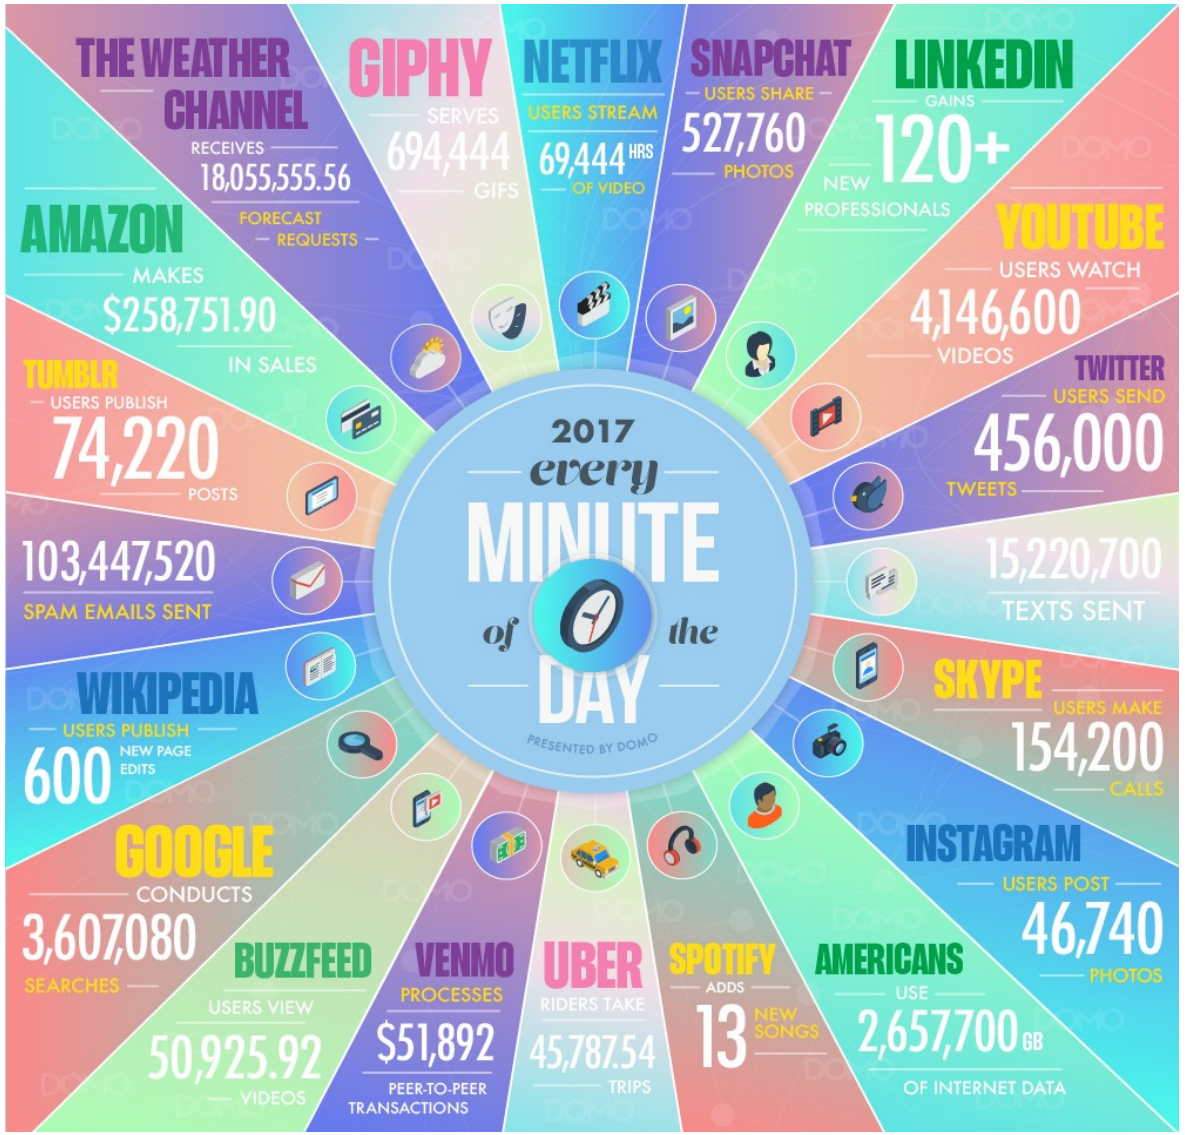

What are the roles in a Big Data organization?

Data Manager

Data engineer

Data Architect

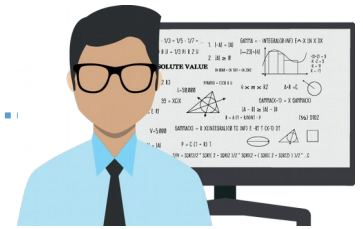

Data Scientist

Machine Learning Engineer

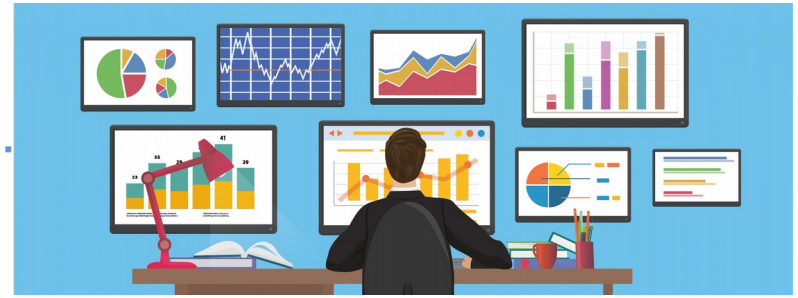

Data Analyst

What is Machine Learning?

Machine learning algorithms build a mathematical model of sample data, known as "training data", in order to make predictions or decisions without being explicitly programmed to perform the task.

Machine learning is used in:

- Email filtering
- Image classification
- Fraud detection
- Etc…

Which use cases we saw at TOTAL?

- Use Case 1: XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
- Use Case 2: XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
- Use Case 3: XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
- Use Case 4: XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX

Some basic vocabulary

Pb: predict the quantity of apples sold in a supermarket for a given day

Variable / feature

Target variable

observation

Dataset

	Temp ext (°C)	Day of week	...	Price/ kg for apples
01/01/2017	-10	3	...	2
02/01/2017	-8	4	...	2,03
03/01/2017	5	5	...	2,04
04/01/2017	6	6	...	2,50
05/01/2017	2	7	...	2,50

Apples sold (kg)
34
37
67
64
33
...
87

What are the branches of machine learning?

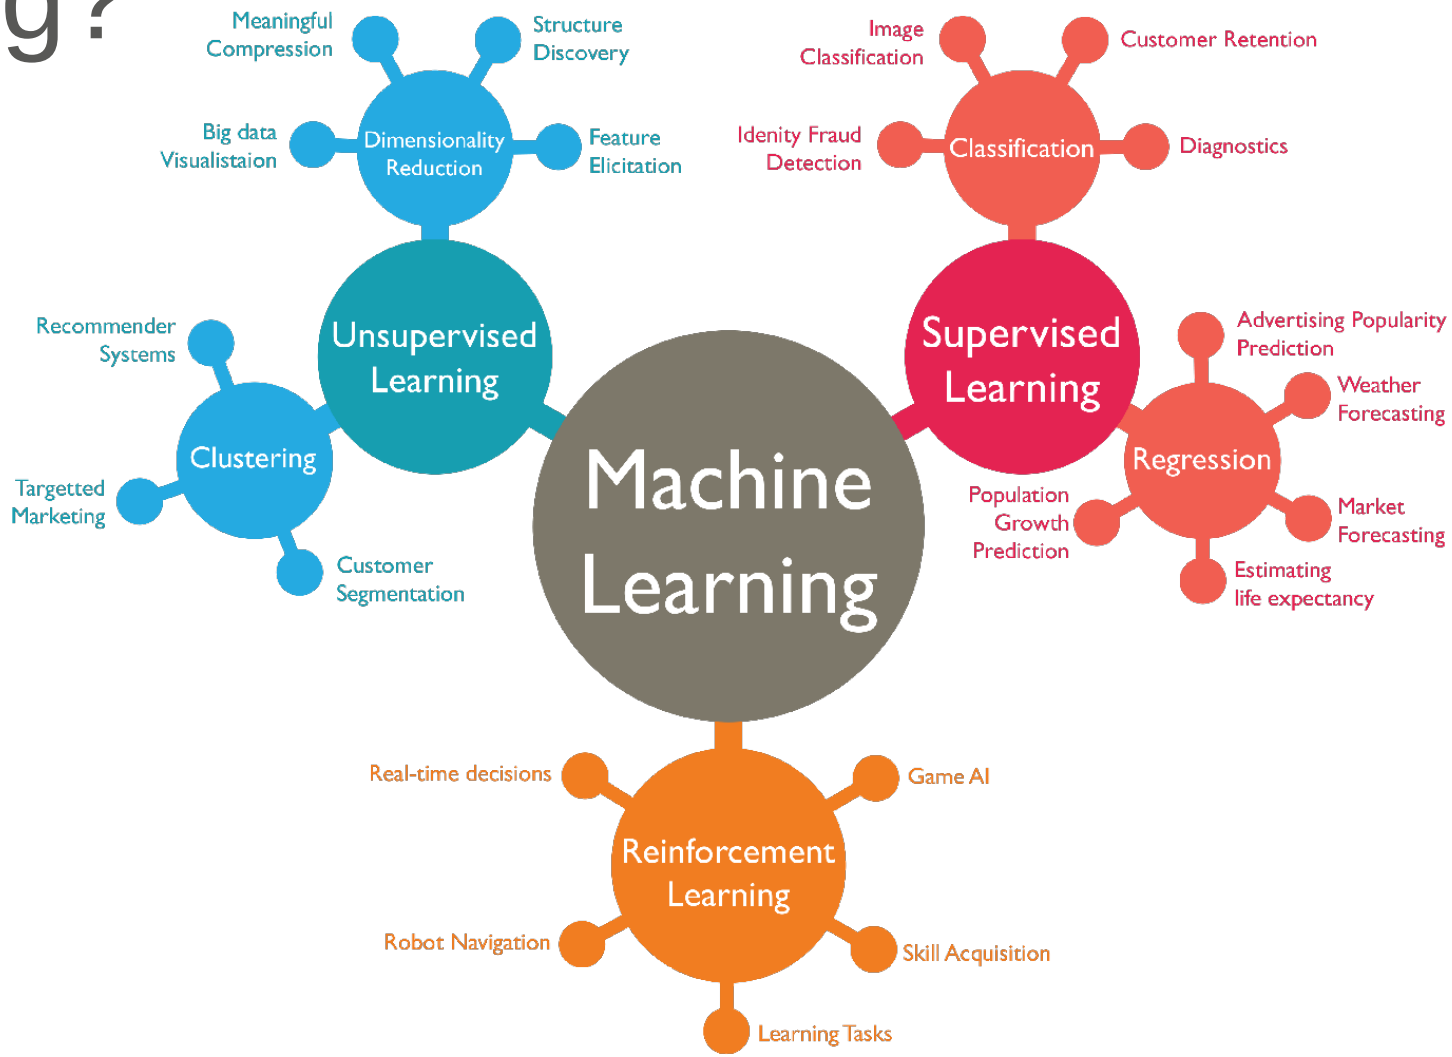

Classification

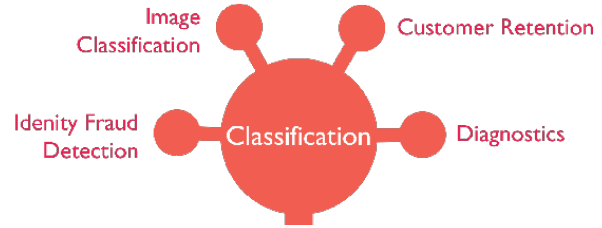

Pb: (Spotify) Will the users buy our premium offer?

	Nb streams per day	Seniority	Buy after trial?
User 1	12	1	⊘
User 2	56	24	○
User 3	467	13	⊘
…	…	…	…
User n	32	4	⊘

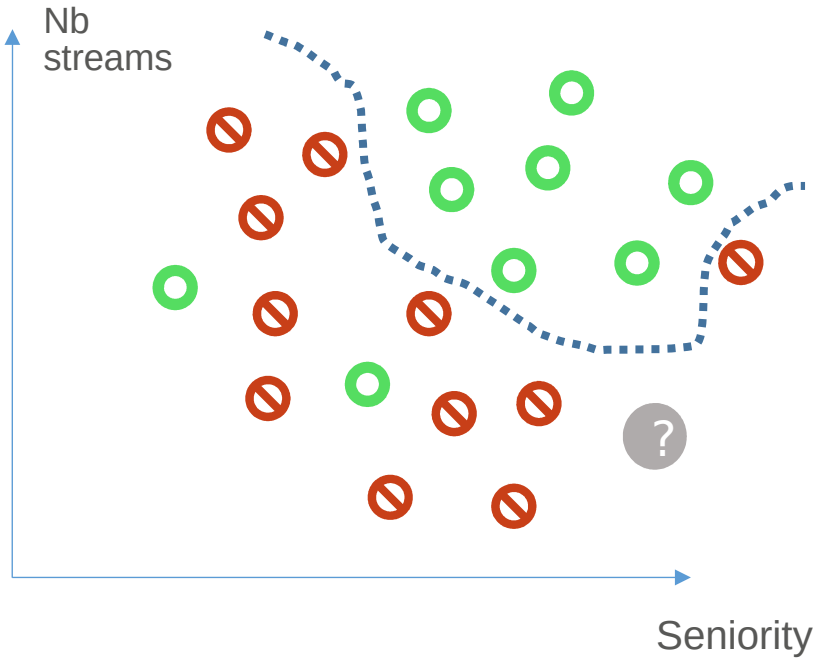

After the model has been fitted to the training set, let's apply the prediction on a new observation:

User n+1			?

The model tries to find the best border that splits the positive and negative observations

Regression

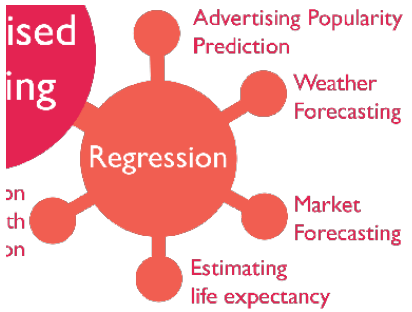

Pb: (MeilleursAgents) How to estimate the price of an appartment?

	Surface (m2)	Price (k€)
Apt 1	12	200 k
Apt 2	56	450 k
Apt 3	130	1200 k
…	…	…
Apt n	32	300 k

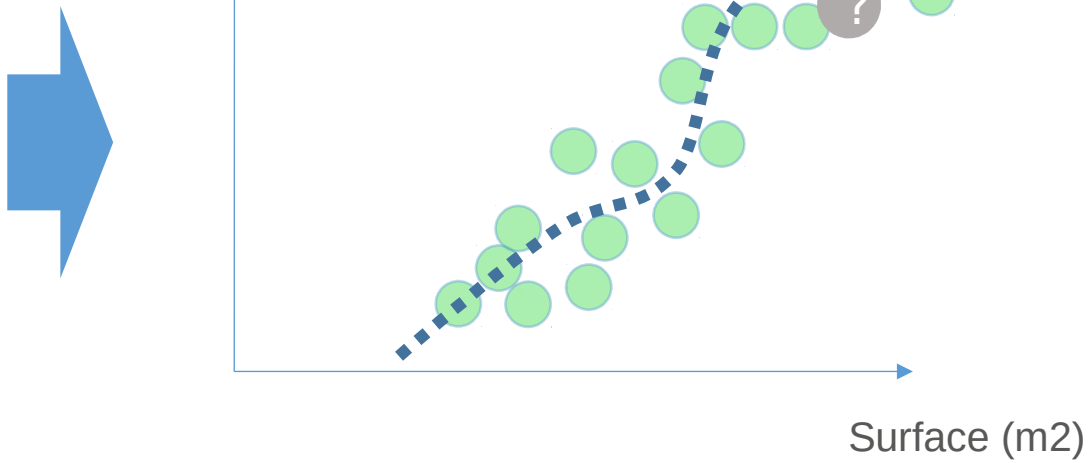

After the model has been fitted to the training set, let's apply the prediction on a new observation:

Apt n+1		?

Clustering

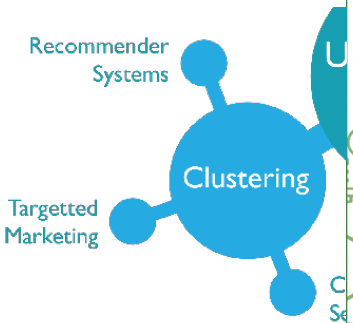

Pb: (Netflix) Can I group similar users by behaviour on the app?

	Nb movies (/month)	Nb connections
User 1	12	1
User 2	32	24
User 3	46	13
…	…	…
User n	32	44

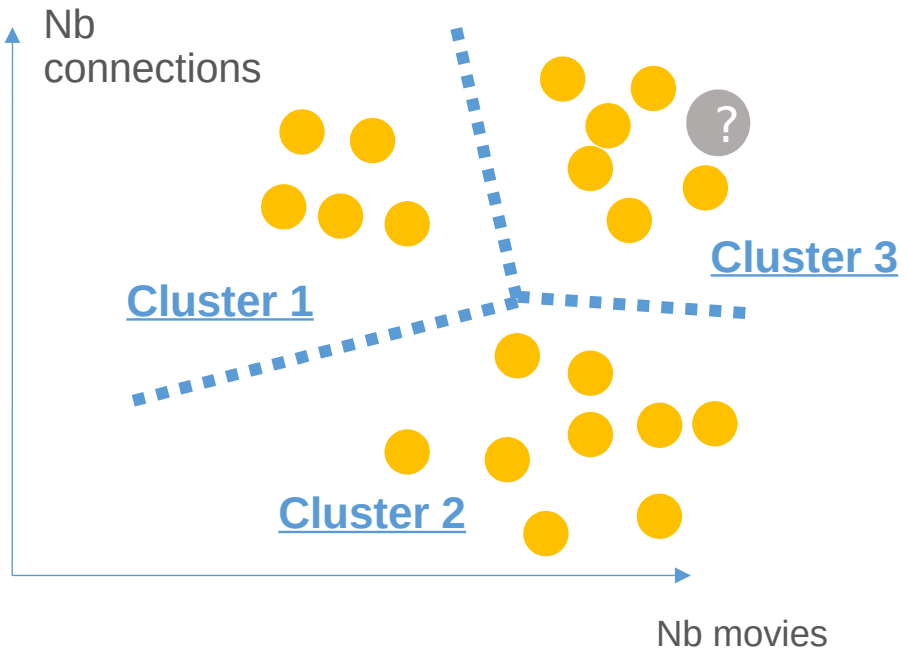

After the model has been fitted to the training set, let's apply the prediction on a new observation:

User n+1			?

The model tries to find the best border that splits the positive and negative observations

Workflow of a ML project

What are we going to do today?

The Challenge

Predict the number of shared bikes rented every hour in San Diego given meteorological information

- **Features**:
  - Temperature
  - Time
  - Humidity
  - Wheather
  - Weekday
  - Is_holiday
  - Etc…
- **Data**: records from 2016/2017
- **Tools**: using Python (via Jupyter Notebook)

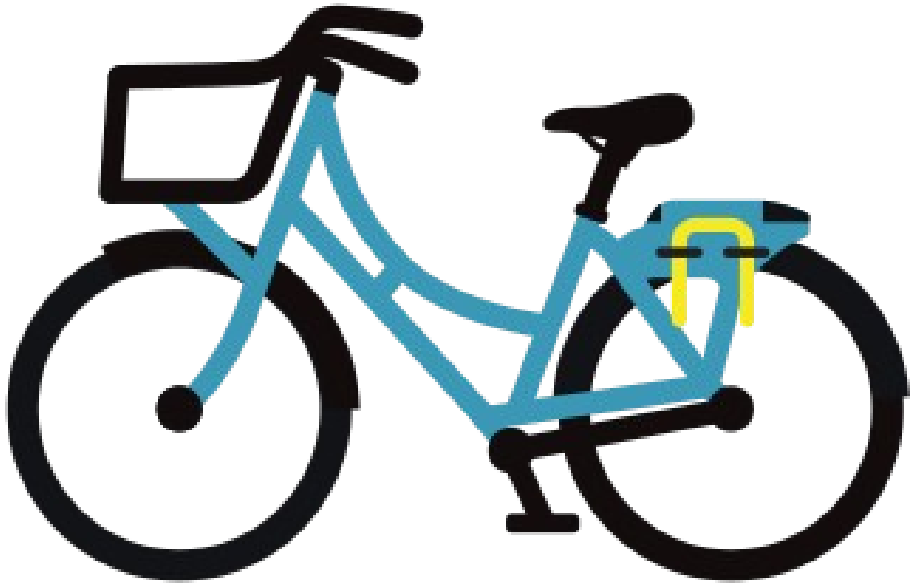

What are we going to do today?

The steps

Theory & Hands on:

- Step 0:
  Introduction to Python and Jupyter
- Step 1:
  Build a first basic model
- Step 2:
  Improve your model: preprocessing
- Step 3:
  Improve your model: model choice and optimise the hyperparameters

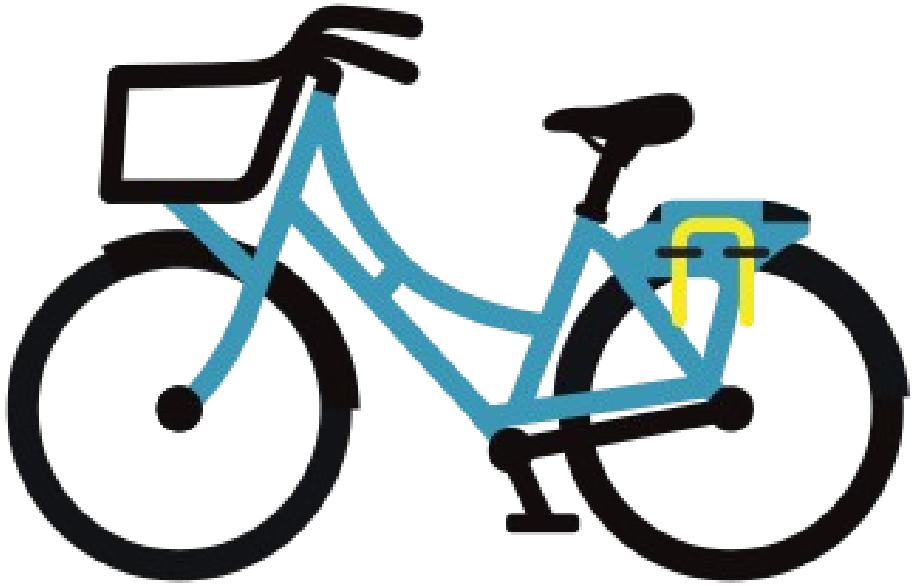

Step 0: Introduction to Python and Jupyter notebooks

Step 1: Build a 1st basic model

- Load data

Do the minimum:

- Remove rows with NA values
- Dummify categorical values
- Split Dataframe in train and test

Do the minimum:

- Choose 1 model only

Do the minimum:

- Choose an evaluation metric
- Compute the score of your model

Step 1: Build a 1st basic model

Split the Dataset in order to evaluate your model

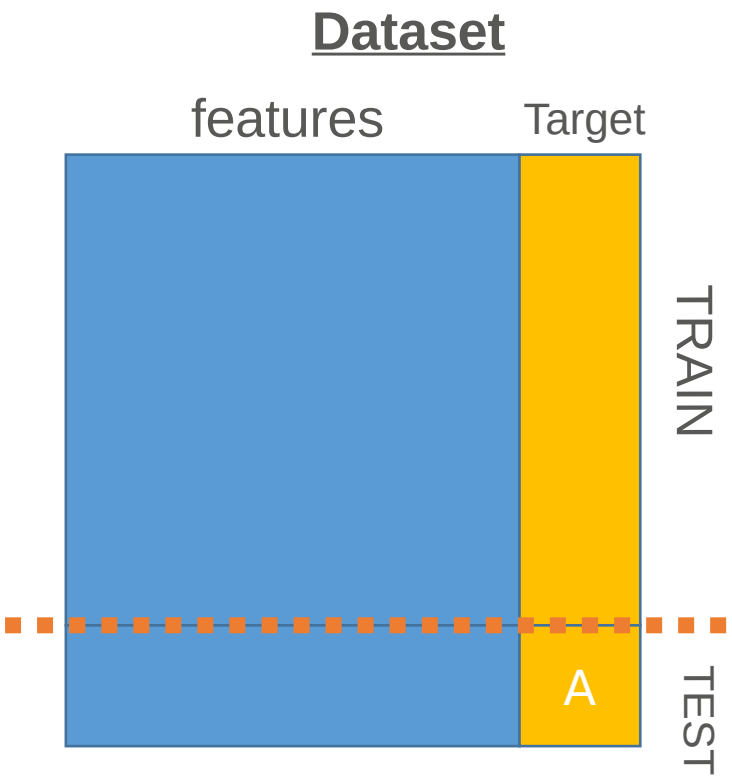

The model trains

Make predictions with the test set

How to compute the performance of a model ?

SCORE

Mae, rmse, etc…

Step 1: Build a 1st basic model

Step 2: Improve your model : Preprocessing

- Load data

- Dummify categorical values
- Split Dataframe in train and test
- Impute missing values
- Add some feature engineering

Do the minimum:

- Choose 1 model only

Do the minimum:

- Choose an evaluation metric
- Compute the score of your model

Step 2: Improve your model : Preprocessing

Step 3: Improve your model : Models and hyperparameters optimisation

Get Data

- Load data

Preprocessing

- Dummify categorical values
- Split Dataframe in train and test
- Impute missing values
- Add some feature engineering

Train the model

- Try several models (NNs, RFs, Gradient Boosting etc…)
- Use several sets of hyperparameters

Evaluate the model

- Choose an evaluation metric
- Compute the score of your model

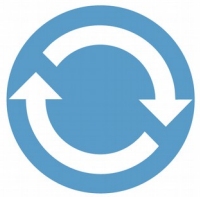

Improve

Step 3: Improve your model : Models and hyperparameters optimisation